

# Parallel Computing with Python

**Monte Lunacek**

Research Computing, University of Colorado **Boulder**

# Outline

- Parallel Computing
  - What is it?
  - Easy parallel example
  - Considerations for using parallelism
- High Throughput Computing IPYthon
  - Configuration
  - Schedules
  - Fault-tolerance
- Conclusion



# Parallel Computing

# What is Parallel Computing?

## Size

- Solve problems that can't fit on a laptop
- Need more than a few GB of RAM
- Need more than a few hundred GB of Disk

## Speed

- Same problem, faster
- Makes a bigger problem more feasible

# What is Supercomputing?



- Definition changes daily!
- Cluster of computers linked together
- 100x bigger, faster, better than a PC

# Janus macbook comparison



Macbook	Janus
2.4 GHz Intel (dual-core)	2.8 GHz Intel (hex-core) X 2 X 1360
3 M cache	12 M cache
8 GB RAM	24 GB RAM

Speed (cpu)                      ~8000X

Size (memory)                      ~4000X

# Parallel Computing

## Shared Memory

- Open MP
- Communication occurs through the shared memory
- Restricted to the number of cores on a node

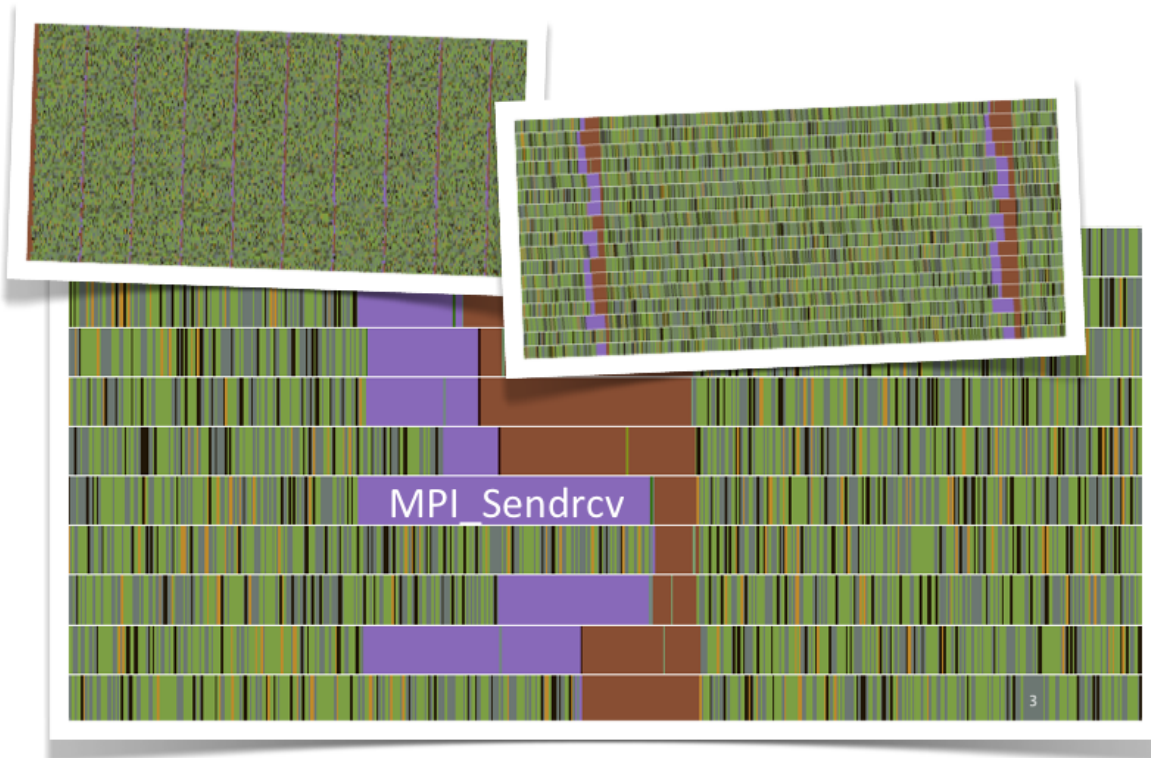

## Distributed

- Message Passing Interface (MPI)
- Communication occurs through message passing
- Use thousands of cores

## Hybrid

- OpenMP + MPI

# Parallel Applications





# High Throughput Computing



## Completely independent tasks

- Simulations
  - Monte Carlo
  - Parameter scan
  - Uncertainty Quantification
  - Parameter Optimization
- Data Analysis (MapReduce)
- Parallel workflows

# Multiprocessing example

## Objectives

- Quick and easy parallelism
- Speedup
- Efficiency
- Karp-Flatt Metric
- Amdahl's law

# IP[y]: Notebook

Notebooks

Clusters

To import a notebook, drag the file onto the listing below or **click here**.

/projects/molu8455/tutorials/python/notebooks

ipython

multiprocessing



# High Throughput Computing with Python

# Success Stories

~500,000 simulations on ~7,000 cores with mpi4py

(<http://mpi4py.scipy.org/>)

Wrapped an engineering simulation with f2py

(<http://www.scipy.org/F2py>) and IPython Parallel

(<http://ipython.org/ipython-doc/dev/parallel/>)

David Folch: parameter optimization with Scoop

(<https://code.google.com/p/scoop/>) and DEAP

(<https://code.google.com/p/deap/>)

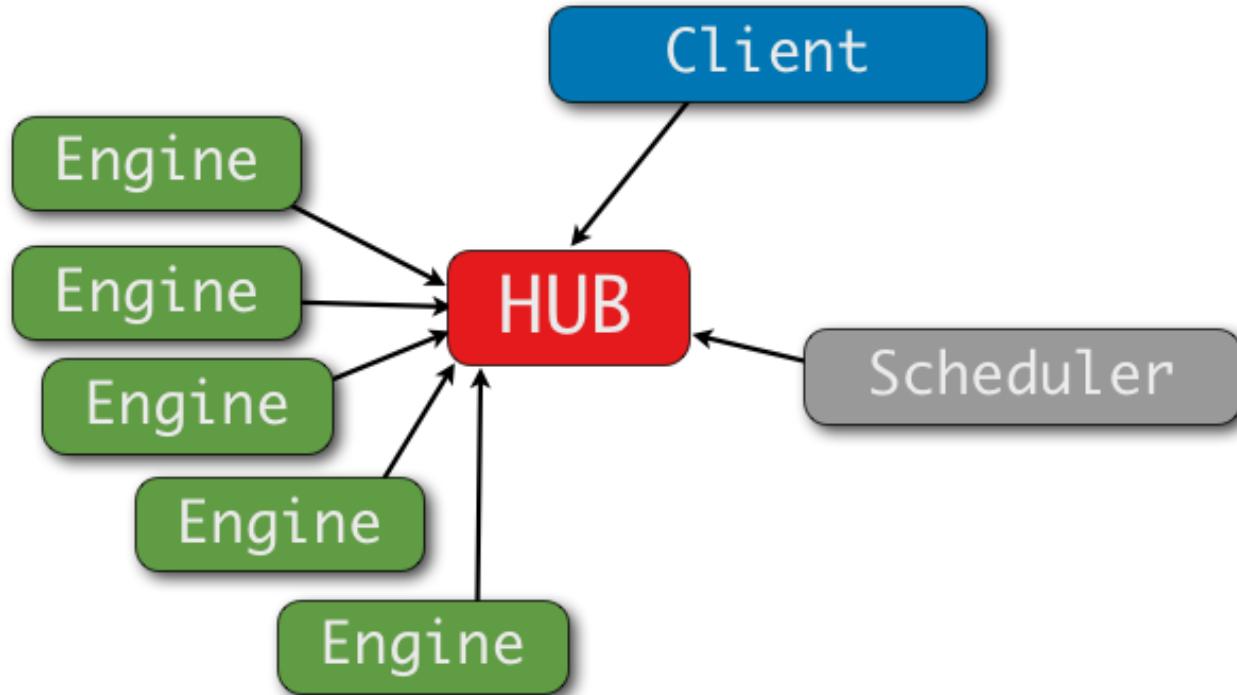
Jon Leff: QIIME IPython Parallel (<http://ipython.org/ipython-doc/dev/parallel/>)

Ann Deml: MPI tasks with Jinja2 (<http://jinja.pocoo.org/docs/>)

# iPython parallel notebook

*Interactive parallel computing on a cluster from a web browser*

# iPython parallel architecture



Goal: launch the engines and the controller client.

# iPython parallel configuration

Launch with SSH and MPI

```
ipython profile create --parallel --profile=ssh
```

```
ipython profile create --parallel --profile=mpi
```

This creates two directories

- `.ipython/profile_ssh` and `.ipython/profile_mpi`

Each directory contains

- `ipython_notebook_config.py`
- `ipcluster_config.py`
- `ipengine_config.py`
- `ipcontroller_config.py`



# iPython parallel configuration

ipython\_notebook\_config.py

```
c = get_config()
c.NotebookApp.ip = '*'
c.NotebookApp.port = 8888
c.NotebookApp.open_browser = False
```

# iPython parallel configuration

profile\_ssh/ipcluster\_config.py

```
c = get_config()
c.LocalControllerLauncher.controller_args = ["--ip='*']
c.IPClusterEngines.engine_launcher_class = 'SSH'
c.SSHEngineLauncher.remote_profile_dir = u'.ipython/profile_ssh'
c.SSHEngineLauncher.ssh_args = ['-t']
```

```
c.IPClusterStart.n = 12

c.SSHEngineSetLauncher.engines = {}
filename = os.environ['PBS_NODEFILE']
with open(filename, 'r') as file:
    for line in file:
        node = line.split()[0]
        if node not in c.SSHEngineSetLauncher.engines:
            c.SSHEngineSetLauncher.engines[node] = c.IPCluster
Start.n
```

# Running iPython parallel

Start the engines

```
ipcluster start --profile='ssh' &
```

Stop the engines

```
ipcluster stop --profile='ssh'
```

# Running the notebook on a cluster

Local

```
ssh -Y molu8455@login.rc.colorado.edu
```

Remote

```
qsub -I -X -q janus-admin -l nodes=2:ppn=12
```

```
ipcluster start --profile='ssh' &
```

```
ipython notebook --profile='ssh' --pylab=inline
```

Local

```
ssh -L 2000:node0379:8888 -f -N molu8455@login.rc.colorado.edu
```

# IP[y]: Notebook

Notebooks

Clusters

To import a notebook, drag the file onto the listing below or **click here**.

/projects/molu8455/tutorials/python/notebooks

ipython

multiprocessing

# IPython Parallel

```
from IPython.parallel import Client
```

Map the values

```
if __name__ == '__main__':  
  
    data = range(200) # tasks  
    rc = Client(profile='ssh')  
    lview = rc.load_balanced_view()  
  
    results = lview.map(simulation, data)  
    results.wait()
```


# iPython parallel schedules

## Config


```
# 'lru', 'weighted', 'pure', 'leastload'  
c.TaskScheduler.scheme_name = 'leastload'  
c.TaskScheduler.hwm = 1
```

### hwm

- 0 Static schedule
- 1 Can have at most one waiting
- 2 Can have at most two waiting










# Tenacious Robustness Test

```
@require('time','socket','random',  
        'IPython.parallel.error.KernelError')  
def simulation(x):  
  
    time.sleep(5)  
    if random.random() < 0.3:  
        raise KernelError  
    return {'task':x,  
            'host' :socket.gethostname()}
```

Launch 10 nodes

Run **several** tasks

At some point, **kill a node**



# Conclusions

Python makes **supercomputing accessible**

Why parallel computing?

- It's the future!
- Puts you ahead in the game
- More efficient

• Jobs • Bash • Static • Dynamic • Fault

